

Investigate_a_Dataset

December 19, 2022

1 Project: Investigate a Dataset - TMDb movie data

1.1 Table of Contents

Introduction

Data Wrangling

Exploratory Data Analysis

Conclusions

Introduction

1.1.1 Dataset Description

We are working with the TMDb movie dataset that contains at least 10.000 values of movie data released all over the years. When we talk about movies details, we refer to the attributes included in the dataset as follows:

- id
- imdb_id
- popularity
- budget
- revenue
- original_title
- cast
- homepage
- director
- tagline
- keywords
- overview
- runtime
- genres
- production_companies
- release_date
- vote_count
- vote_average
- release_year
- budget_adj
- revenue_adj

Looking at these attributes, it could be important to give some brief explanation of some of them, such as: * `vote_count`: stands for number of votes given to a movie * `vote_average`: stands for the average of vote classification given from the vote count * `budget_adj`: stands for the budget of the associated movie in terms of 2010 dollars, accounting for inflation over time. * `revenue_adj`: stands for the reveue of the associated movie in terms of 2010 dollars, accounting for inflation over time.

1.1.2 Question(s) for Analysis

The present report is planned to answer the following questions during the analysis course:

1. Which genre of movie tends to be more appreciated (in terms of voting)?
2. Has nowadays movies been well received by people comparing to old movies?
3. What is the best month of the year used to release the movie?
4. Ranking of the best production companies from year to year on TMDb. By how much?
5. Does long runtime movies are more expensive to produce compared to short runtime movies?
6. Which genre of movie trends to be more expensive to produce?
7. Which movie director directed the production of movies (more than one) that became popular in relation to others, in the most recent year in the dataset?

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
% matplotlib inline
```

```
In [2]: # Upgrade pandas to use dataframe.explode() function.
!pip install --upgrade pandas
```

```
Requirement already up-to-date: pandas in /opt/conda/lib/python3.6/site-packages (1.1.5)
Requirement already satisfied, skipping upgrade: numpy>=1.15.4 in /opt/conda/lib/python3.6/site-
Requirement already satisfied, skipping upgrade: python-dateutil>=2.7.3 in /opt/conda/lib/python
Requirement already satisfied, skipping upgrade: pytz>=2017.2 in /opt/conda/lib/python3.6/site-p
Requirement already satisfied, skipping upgrade: six>=1.5 in /opt/conda/lib/python3.6/site-packa
```

Data Wrangling

1.1.3 General Properties

For this process, it is necessary to explore the general properties of the dataset, the way that it can lead us to decide what actions can be done to turn the dataset into a clean and accurate one for analysis.

```
In [3]: df = pd.read_csv('tmdb-movies.csv')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	id	imdb_id	popularity	budget	revenue	\
0	135397	tt0369610	32.985763	150000000	1513528810	
1	76341	tt1392190	28.419936	150000000	378436354	
2	262500	tt2908446	13.112507	110000000	295238201	
3	140607	tt2488496	11.173104	200000000	2068178225	
4	168259	tt2820852	9.335014	190000000	1506249360	

	original_title	\
0	Jurassic World	
1	Mad Max: Fury Road	
2	Insurgent	
3	Star Wars: The Force Awakens	
4	Furious 7	

	cast	\
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	
2	Shailene Woodley Theo James Kate Winslet Ansel...	
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...	
4	Vin Diesel Paul Walker Jason Statham Michelle ...	

	homepage	director	\
0	http://www.jurassicworld.com/	Colin Trevorrow	
1	http://www.madmaxmovie.com/	George Miller	
2	http://www.thedivergentseries.movie/#insurgent	Robert Schwentke	
3	http://www.starwars.com/films/star-wars-episod...	J.J. Abrams	
4	http://www.furious7.com/	James Wan	

	tagline	...	\
0	The park is open.	...	
1	What a Lovely Day.	...	
2	One Choice Can Destroy You	...	
3	Every generation has a story.	...	
4	Vengeance Hits Home	...	

	overview	runtime	\
0	Twenty-two years after the events of Jurassic ...	124	
1	An apocalyptic story set in the furthest reach...	120	
2	Beatrice Prior must confront her inner demons ...	119	
3	Thirty years after defeating the Galactic Empi...	136	
4	Deckard Shaw seeks revenge against Dominic Tor...	137	

	genres	\
0	Action Adventure Science Fiction Thriller	
1	Action Adventure Science Fiction Thriller	
2	Adventure Science Fiction Thriller	

```

3 Action|Adventure|Science Fiction|Fantasy
4 Action|Crime|Thriller

```

```

                                production_companies release_date vote_count \
0 Universal Studios|Amblin Entertainment|Legenda...      6/9/15      5562
1 Village Roadshow Pictures|Kennedy Miller Produ...      5/13/15      6185
2 Summit Entertainment|Mandeville Films|Red Wago...      3/18/15      2480
3 Lucasfilm|Truenorth Productions|Bad Robot        12/15/15      5292
4 Universal Pictures|Original Film|Media Rights ...      4/1/15      2947

```

```

    vote_average release_year budget_adj revenue_adj
0           6.5         2015  1.379999e+08  1.392446e+09
1           7.1         2015  1.379999e+08  3.481613e+08
2           6.3         2015  1.012000e+08  2.716190e+08
3           7.5         2015  1.839999e+08  1.902723e+09
4           7.3         2015  1.747999e+08  1.385749e+09

```

[5 rows x 21 columns]

- Show the dimentions of the dataset.

```
In [5]: df.shape
```

```
Out[5]: (10866, 21)
```

- Show the data type of each column in the dataset.

```
In [6]: df.dtypes
```

```

Out[6]: id                int64
imdb_id                  object
popularity              float64
budget                  int64
revenue                 int64
original_title          object
cast                   object
homepage                object
director                object
tagline                 object
keywords                object
overview                object
runtime                 int64
genres                  object
production_companies    object
release_date            object
vote_count              int64
vote_average            float64
release_year            int64
budget_adj              float64

```

```
revenue_adj          float64
dtype: object
```

- Show the number of not null values in each column, so as their data types, in the dataset.

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    10866 non-null  int64
1   imdb_id              10856 non-null  object
2   popularity           10866 non-null  float64
3   budget               10866 non-null  int64
4   revenue              10866 non-null  int64
5   original_title       10866 non-null  object
6   cast                 10790 non-null  object
7   homepage             2936 non-null   object
8   director             10822 non-null  object
9   tagline              8042 non-null   object
10  keywords              9373 non-null   object
11  overview             10862 non-null  object
12  runtime              10866 non-null  int64
13  genres               10843 non-null  object
14  production_companies  9836 non-null   object
15  release_date         10866 non-null  object
16  vote_count           10866 non-null  int64
17  vote_average         10866 non-null  float64
18  release_year         10866 non-null  int64
19  budget_adj           10866 non-null  float64
20  revenue_adj          10866 non-null  float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

- Show the number of null values in each column.

```
In [8]: df.isnull().sum()
```

```
Out[8]: id                    0
        imdb_id              10
        popularity           0
        budget               0
        revenue              0
        original_title       0
        cast                 76
        homepage            7930
```

director	44
tagline	2824
keywords	1493
overview	4
runtime	0
genres	23
production_companies	1030
release_date	0
vote_count	0
vote_average	0
release_year	0
budget_adj	0
revenue_adj	0
dtype:	int64

- Show how many duplicated rows are in the dataset.

```
In [9]: sum(df.duplicated())
```

```
Out[9]: 1
```

- Show how many unique values are in each column in the dataset.

```
In [10]: df.nunique()
```

```
Out[10]: id                10865
imdb_id                  10855
popularity               10814
budget                   557
revenue                  4702
original_title          10571
cast                    10719
homepage                 2896
director                 5067
tagline                  7997
keywords                 8804
overview                10847
runtime                  247
genres                  2039
production_companies     7445
release_date             5909
vote_count              1289
vote_average             72
release_year             56
budget_adj              2614
revenue_adj             4840
dtype: int64
```

- Show some statistic results of all numerical columns in the dataset.

```
In [11]: df.describe()
```

```
Out[11]:
```

	id	popularity	budget	revenue	runtime \
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000

	vote_count	vote_average	release_year	budget_adj	revenue_adj
count	10866.000000	10866.000000	10866.000000	1.086600e+04	1.086600e+04
mean	217.389748	5.974922	2001.322658	1.755104e+07	5.136436e+07
std	575.619058	0.935142	12.812941	3.430616e+07	1.446325e+08
min	10.000000	1.500000	1960.000000	0.000000e+00	0.000000e+00
25%	17.000000	5.400000	1995.000000	0.000000e+00	0.000000e+00
50%	38.000000	6.000000	2006.000000	0.000000e+00	0.000000e+00
75%	145.750000	6.600000	2011.000000	2.085325e+07	3.369710e+07
max	9767.000000	9.200000	2015.000000	4.250000e+08	2.827124e+09

1.1.4 Data Cleaning

From the investigation of the data above and based on the questions to analyse, the following cleaning steps are going to be taken:

- Budget : transform to float
- Revenue : transform to float

```
In [12]: df['budget'] = df['budget'].astype(float) # Change the budget data type to float
```

```
In [13]: df['revenue'] = df['revenue'].astype(float) # Change the revenue data type to float
```

- Release date: transform to date

```
In [14]: def arrange_date(x):
    if (int(x[2]) > 20):
        x[2] = x[2].replace(x[2], '19'+x[2])
    else:
        x[2] = x[2].replace(x[2], '20'+x[2])
    if (int(x[1]) < 10):
        x[1] = x[1].replace(x[1], '0'+x[1])
    if (int(x[0]) < 10):
        x[0] = x[0].replace(x[0], '0'+x[0])
    x = x[0]+x[1]+x[2]
    return x
```

```
In [15]: df['release_date'] = df['release_date'].apply(lambda x: x.split('/')) # Calls a function
df['release_date'] = df['release_date'].apply(lambda x: arrange_date(x)) # Calls the function
df['release_date'] = pd.to_datetime(df['release_date'], format='%m/%d/%Y') #Format the date
```

- There is one duplicated arrow, so it can be dropped as follows.

```
In [16]: df = df.drop_duplicates()

sum(df.duplicated())
```

```
Out[16]: 0
```

- There are various null rows in some columns of String data type, but there is no need to fill any value in them, due to the description of each of column, turning difficult to preview what type of value could fill the columns. Fortunately, those columns are "Tagline", "Keywords", "Overview" and "homepage" that are not needed for the analysis. So they can be dropped as follows.

```
In [17]: df.drop(['tagline', 'keywords', 'overview', 'homepage'], axis = 1, inplace = True) # drop
```

- Beside these columns, there are columns like cast, genre and production_companies that will be used to answer the questions. Therefore, having null values in them will not help to make the analysis. So, the null values in these columns will be dropped so as the rows that contains these null values.

```
In [18]: df.dropna(inplace = True) # drop rows with any null values in the dataset
df.isnull().sum()
```

```
Out[18]: id                0
imdb_id                  0
popularity               0
budget                  0
revenue                 0
original_title           0
cast                    0
director                0
runtime                 0
genres                  0
production_companies     0
release_date            0
vote_count              0
vote_average            0
release_year            0
budget_adj              0
revenue_adj             0
dtype: int64
```


- Genres and Production_Companies columns have multiple values that were separated with '|' in each row. turning difficult to make some of the analisys needed. It will be necessary to trim each column value as follows.
- Based in the questions defined above, it will be created two separated datasets (that contains the rows with the '|' character in 'genre' and 'production_companies' columns). One of the dataframes will contain the 'genres' column trimmed, and other the 'production_companies' column trimmed for later analysis.

```
In [19]: g_df = df[df['genres'].str.contains('|')] # dataframe containing multiple genre values
        p_df = df[df['production_companies'].str.contains('|')] # dataframe containing multiple
```

```
In [20]: print(g_df.shape)
        print(p_df.shape)
```

```
(9770, 17)
```

```
(9770, 17)
```

- As the results above, it is clear that all the columns contains 9770 rows as the cleaned dataframe.

- There's a function that will help us trim and separate each value in the rows of a determined column.

```
In [21]: def split_values(x, i, df, column):
        ind = df[df[column] == x].index.values # gets the index of the row
        x = x.split('|') # split the multiple values in each row of the column
        if (i == 0):
            return x[i]
        elif (i < len(x)):
            return x[i] # check if the row has any value in the index i after splitting
        else:
            df.drop(ind) # drop the row that contain values trimmed in other dataframe in
```

- For separating values in 'genre' column, it will be needed hybrid dataframes. We will use the

```
In [22]: g1_df = g_df.copy()
        g2_df = g_df.copy()
        g1_df['genres'] = g_df['genres'].apply(lambda x: split_values(x, 0, g_df, 'genres'))
        g2_df['genres'] = g_df['genres'].apply(lambda x: split_values(x, 1, g_df, 'genres'))
```

- Having the two hybrid dataframes with the first and second value of each row in genre column,

```
In [23]: g_df = pd.concat([g1_df, g2_df])
        g_df
```

```
Out[23]:
```

	id	imdb_id	popularity	budget	revenue	\
0	135397	tt0369610	32.985763	150000000.0	1.513529e+09	
1	76341	tt1392190	28.419936	150000000.0	3.784364e+08	
2	262500	tt2908446	13.112507	110000000.0	2.952382e+08	
3	140607	tt2488496	11.173104	200000000.0	2.068178e+09	

4	168259	tt2820852	9.335014	1900000000.0	1.506249e+09
...
10861	21	tt0060371	0.080598	0.0	0.000000e+00
10862	20379	tt0060472	0.065543	0.0	0.000000e+00
10863	39768	tt0060161	0.065141	0.0	0.000000e+00
10864	21449	tt0061177	0.064317	0.0	0.000000e+00
10865	22293	tt0060666	0.035919	19000.0	0.000000e+00

	original_title	\
0	Jurassic World	
1	Mad Max: Fury Road	
2	Insurgent	
3	Star Wars: The Force Awakens	
4	Furious 7	
...	...	
10861	The Endless Summer	
10862	Grand Prix	
10863	Beregis Avtomobilya	
10864	What's Up, Tiger Lily?	
10865	Manos: The Hands of Fate	

	cast	director	\
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow	
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	George Miller	
2	Shailene Woodley Theo James Kate Winslet Ansel...	Robert Schwentke	
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...	J.J. Abrams	
4	Vin Diesel Paul Walker Jason Statham Michelle ...	James Wan	
...	
10861	Michael Hynson Robert August Lord 'Tally Ho' B...	Bruce Brown	
10862	James Garner Eva Marie Saint Yves Montand Tosh...	John Frankenheimer	
10863	Innokentiy Smoktunovskiy Oleg Efremov Georgi Z...	Eldar Ryazanov	
10864	Tatsuya Mihashi Akiko Wakabayashi Mie Hama Joh...	Woody Allen	
10865	Harold P. Warren Tom Neyman John Reynolds Dian...	Harold P. Warren	

	runtime	genres	production_companies	\
0	124	Action	Universal Studios Amblin Entertainment Legenda...	
1	120	Action	Village Roadshow Pictures Kennedy Miller Produ...	
2	119	Adventure	Summit Entertainment Mandeville Films Red Wago...	
3	136	Action	Lucasfilm Truenorth Productions Bad Robot	
4	137	Action	Universal Pictures Original Film Media Rights ...	
...	
10861	95	None	Bruce Brown Films	
10862	176	Adventure	Cherokee Productions Joel Productions Douglas ...	
10863	94	Comedy	Mosfilm	
10864	80	Comedy	Benedict Pictures Corp.	
10865	74	None	Norm-Iris	

release_date	vote_count	vote_average	release_year	budget_adj	\
--------------	------------	--------------	--------------	------------	---

0	2015-06-09	5562	6.5	2015	1.379999e+08
1	2015-05-13	6185	7.1	2015	1.379999e+08
2	2015-03-18	2480	6.3	2015	1.012000e+08
3	2015-12-15	5292	7.5	2015	1.839999e+08
4	2015-04-01	2947	7.3	2015	1.747999e+08
...
10861	1966-06-15	11	7.4	1966	0.000000e+00
10862	1966-12-21	20	5.7	1966	0.000000e+00
10863	1966-01-01	11	6.5	1966	0.000000e+00
10864	1966-11-02	22	5.4	1966	0.000000e+00
10865	1966-11-15	15	1.5	1966	1.276423e+05

	revenue_adj
0	1.392446e+09
1	3.481613e+08
2	2.716190e+08
3	1.902723e+09
4	1.385749e+09
...	...
10861	0.000000e+00
10862	0.000000e+00
10863	0.000000e+00
10864	0.000000e+00
10865	0.000000e+00

[19540 rows x 17 columns]

- For each 'production_companies' row we will use in maximum first two values as sample:

```
In [24]: p1_df = p_df.copy()
p2_df = p_df.copy()
p1_df['production_companies'] = p_df['production_companies'].apply(lambda x: split_valu
p2_df['production_companies'] = p_df['production_companies'].apply(lambda x: split_valu
```

- Finally, the hybrid dataframes will be joined as follows.

```
In [25]: p_df = pd.concat([p1_df, p2_df])
p_df
```

```
Out[25]:
```

	id	imdb_id	popularity	budget	revenue \
0	135397	tt0369610	32.985763	150000000.0	1.513529e+09
1	76341	tt1392190	28.419936	150000000.0	3.784364e+08
2	262500	tt2908446	13.112507	110000000.0	2.952382e+08
3	140607	tt2488496	11.173104	200000000.0	2.068178e+09
4	168259	tt2820852	9.335014	190000000.0	1.506249e+09
...
10861	21	tt0060371	0.080598	0.0	0.000000e+00
10862	20379	tt0060472	0.065543	0.0	0.000000e+00
10863	39768	tt0060161	0.065141	0.0	0.000000e+00

10864	21449	tt0061177	0.064317	0.0	0.000000e+00
10865	22293	tt0060666	0.035919	19000.0	0.000000e+00

	original_title \
0	Jurassic World
1	Mad Max: Fury Road
2	Insurgent
3	Star Wars: The Force Awakens
4	Furious 7
...	...
10861	The Endless Summer
10862	Grand Prix
10863	Beregis Avtomobilya
10864	What's Up, Tiger Lily?
10865	Manos: The Hands of Fate

	cast	director \
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	Colin Trevorrow
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	George Miller
2	Shailene Woodley Theo James Kate Winslet Ansel...	Robert Schwentke
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...	J.J. Abrams
4	Vin Diesel Paul Walker Jason Statham Michelle ...	James Wan
...
10861	Michael Hynson Robert August Lord 'Tally Ho' B...	Bruce Brown
10862	James Garner Eva Marie Saint Yves Montand Tosh...	John Frankenheimer
10863	Innokentiy Smoktunovskiy Oleg Efremov Georgi Z...	Eldar Ryazanov
10864	Tatsuya Mihashi Akiko Wakabayashi Mie Hama Joh...	Woody Allen
10865	Harold P. Warren Tom Neyman John Reynolds Dian...	Harold P. Warren

	runtime	genres \
0	124	Action Adventure Science Fiction Thriller
1	120	Action Adventure Science Fiction Thriller
2	119	Adventure Science Fiction Thriller
3	136	Action Adventure Science Fiction Fantasy
4	137	Action Crime Thriller
...
10861	95	Documentary
10862	176	Action Adventure Drama
10863	94	Mystery Comedy
10864	80	Action Comedy
10865	74	Horror

	production_companies	release_date	vote_count	vote_average \
0	Universal Studios	2015-06-09	5562	6.5
1	Village Roadshow Pictures	2015-05-13	6185	7.1
2	Summit Entertainment	2015-03-18	2480	6.3
3	Lucasfilm	2015-12-15	5292	7.5
4	Universal Pictures	2015-04-01	2947	7.3

...
10861	None	1966-06-15	11	7.4
10862	Joel Productions	1966-12-21	20	5.7
10863	None	1966-01-01	11	6.5
10864	None	1966-11-02	22	5.4
10865	None	1966-11-15	15	1.5

	release_year	budget_adj	revenue_adj
0	2015	1.379999e+08	1.392446e+09
1	2015	1.379999e+08	3.481613e+08
2	2015	1.012000e+08	2.716190e+08
3	2015	1.839999e+08	1.902723e+09
4	2015	1.747999e+08	1.385749e+09
...
10861	1966	0.000000e+00	0.000000e+00
10862	1966	0.000000e+00	0.000000e+00
10863	1966	0.000000e+00	0.000000e+00
10864	1966	0.000000e+00	0.000000e+00
10865	1966	1.276423e+05	0.000000e+00

[19540 rows x 17 columns]

- It was noted that 'budget' and 'revenue' columns do not have null values but have values represented as '0.0', that can be described as null. It is show below:

```
In [26]: df[df['revenue'] == 0.0]
```

```
Out[26]:
```

	id	imdb_id	popularity	budget	revenue	\
48	265208	tt2231253	2.932340	30000000.0	0.0	
67	334074	tt3247714	2.331636	20000000.0	0.0	
74	347096	tt3478232	2.165433	0.0	0.0	
75	308369	tt2582496	2.141506	0.0	0.0	
92	370687	tt3608646	1.876037	0.0	0.0	
...
10861	21	tt0060371	0.080598	0.0	0.0	
10862	20379	tt0060472	0.065543	0.0	0.0	
10863	39768	tt0060161	0.065141	0.0	0.0	
10864	21449	tt0061177	0.064317	0.0	0.0	
10865	22293	tt0060666	0.035919	19000.0	0.0	
			original_title	\		
48			Wild Card			
67			Survivor			
74			Mythica: The Darkspore			
75			Me and Earl and the Dying Girl			
92			Mythica: The Necromancer			
...			...			
10861			The Endless Summer			

10862	Grand Prix
10863	Beregis Avtomobilya
10864	What's Up, Tiger Lily?
10865	Manos: The Hands of Fate

	cast	director \
48	Jason Statham Michael Angarano Milo Ventimigli...	Simon West
67	Pierce Brosnan Milla Jovovich Dylan McDermott ...	James McTeigue
74	Melanie Stone Kevin Sorbo Adam Johnson Jake St...	Anne K. Black
75	Thomas Mann RJ Cyler Olivia Cooke Connie Britt...	Alfonso Gomez-Rejon
92	Melanie Stone Adam Johnson Kevin Sorbo Nicola ...	A. Todd Smith
...
10861	Michael Hynson Robert August Lord 'Tally Ho' B...	Bruce Brown
10862	James Garner Eva Marie Saint Yves Montand Tosh...	John Frankenheimer
10863	Innokentiy Smoktunovskiy Oleg Efremov Georgi Z...	Eldar Ryazanov
10864	Tatsuya Mihashi Akiko Wakabayashi Mie Hama Joh...	Woody Allen
10865	Harold P. Warren Tom Neyman John Reynolds Dian...	Harold P. Warren

	runtime	genres \
48	92	Thriller Crime Drama
67	96	Crime Thriller Action
74	108	Action Adventure Fantasy
75	105	Comedy Drama
92	0	Fantasy Action Adventure
...
10861	95	Documentary
10862	176	Action Adventure Drama
10863	94	Mystery Comedy
10864	80	Action Comedy
10865	74	Horror

	production_companies	release_date \
48	Current Entertainment Lionsgate Sierra / Affin...	2015-01-14
67	Nu Image Films Winkler Films Millennium Films ...	2015-05-21
74	Arrowstorm Entertainment	2015-06-24
75	Indian Paintbrush	2015-06-12
92	Arrowstorm Entertainment Camera 40 Productions...	2015-12-19
...
10861	Bruce Brown Films	1966-06-15
10862	Cherokee Productions Joel Productions Douglas ...	1966-12-21
10863	Mosfilm	1966-01-01
10864	Benedict Pictures Corp.	1966-11-02
10865	Norm-Iris	1966-11-15

	vote_count	vote_average	release_year	budget_adj	revenue_adj
48	481	5.3	2015	2.759999e+07	0.0
67	280	5.4	2015	1.839999e+07	0.0
74	27	5.1	2015	0.000000e+00	0.0

75	569	7.7	2015	0.000000e+00	0.0
92	11	5.4	2015	0.000000e+00	0.0
...
10861	11	7.4	1966	0.000000e+00	0.0
10862	20	5.7	1966	0.000000e+00	0.0
10863	11	6.5	1966	0.000000e+00	0.0
10864	22	5.4	1966	0.000000e+00	0.0
10865	15	1.5	1966	1.276423e+05	0.0

[5020 rows x 17 columns]

- The number of rows containing these values are more than the half of total rows of the cleaned

Exploratory Data Analysis


For the research of the questions posed in the Introduction section, it will be computed the relevant statistics to compare and show trends in the visualizations related to the data, based in the questions.

1.1.5 1. Which genre of movie tends to be more appreciated (in terms of voting)?

- We will start the research of this question plotting the histogram for Vote Average distribution.

```
In [27]: df['vote_average'].hist(figsize = (3, 3))
         df['vote_average'].median()
```

Out[27]: 6.0



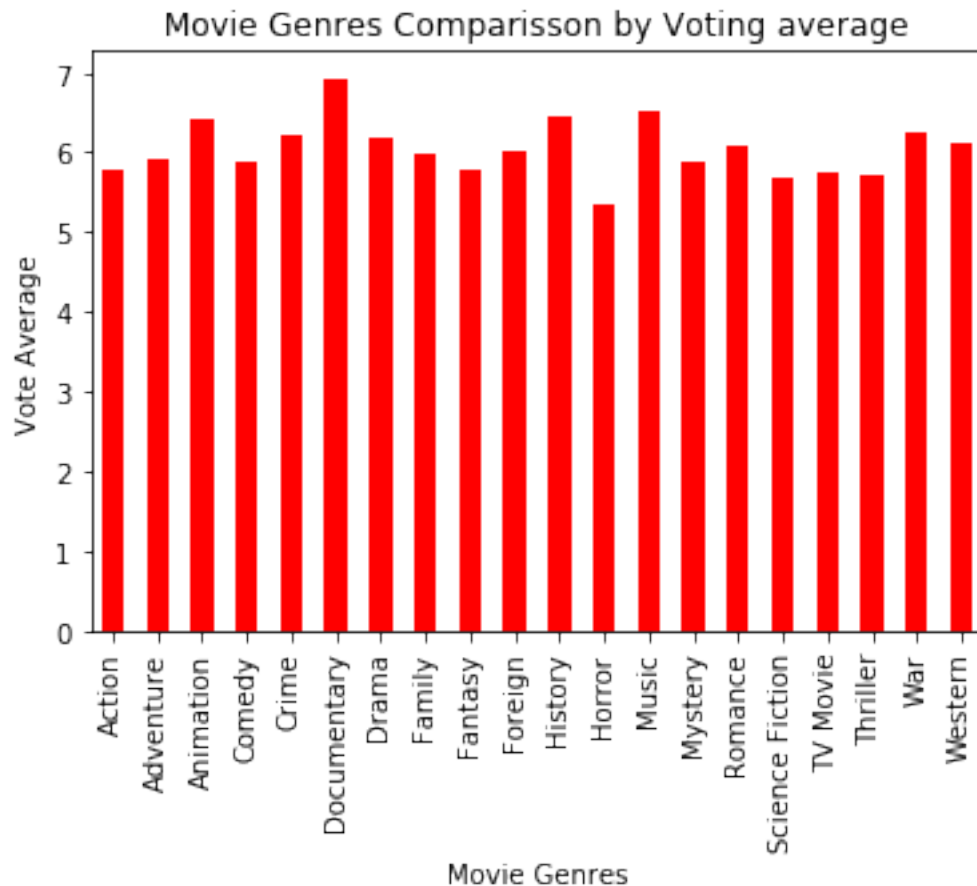
As it is shown above, most of the films tend to have a vote average between 5 to 7 points, and very few of them tend to have more than 8 or less than 4 points as vote.

This plot can lead to the conclusion that the vote average is evenly distributed and not up the sides, stretching across the entire graph. Furthermore, it is balanced toward the center of the frame, with no obvious skew.

- Furthermore, it will be calculated the mean value of vote average in function of each genre of movie using the genre dataframe created, as shown above:

```
In [28]: genre = g_df.groupby(['genres'])['vote_average'].mean()
genre.plot(kind = 'bar', color = 'red', figsize=(6, 4))
plt.title('Movie Genres Comparisson by Voting average')
plt.xlabel('Movie Genres')
plt.ylabel('Vote Average')
```

```
Out[28]: Text(0,0.5,'Vote Average')
```



The plot results tells that Documentary genre could be the most appreciated by people.


Note: Each movie can have more than one genre and for this analysis where used in maximum only two of those values as sample, because of the number of processes that could be executed in large amount of data generated in splitting phase. With that in mind, the statistical result for this question can be different if we decide to add all the data provided by the dataset.

1.1.6 2. Has nowadays movies been well received by people comparing to old movies?

- We will start the research of this question plotting the histogram for Release year distribution of movies.


```
In [29]: df['release_year'].hist(figsize = (5, 3))
```

```
Out[29]: <matplotlib.axes._subplots.AxesSubplot at 0x7efc85ec5da0>
```




As we can see, the number of movies released every year tends to be larger than the previous year.

- Plotting the Scatterplot of Release Year in Relation to Vote Average.

```
In [30]: df.plot(kind = 'scatter', x = 'release_year', y = 'vote_average', figsize = (5, 5))
```

```
Out[30]: <matplotlib.axes._subplots.AxesSubplot at 0x7efc859e2278>
```



This data don't have any kind of pattern, which means that no relationship exists between Release year and Vote average.

- Lastly, for research using the bar plotting, we will define a median first, which function is to separate old movies from modern movies based in the released date. Next step is to define the average vote_average of these two types to compare.


```
In [31]: m = pd.Timestamp(df['release_date'].astype(np.int64).median())
         old_movies = df[df['release_date'] < m]
         old_movies = old_movies['vote_average'].mean()
         modern_movies = df[df['release_date'] >= m]
         modern_movies = modern_movies['vote_average'].mean()
```

Having old movies and modern movies well separated and defined, the results are shown in the plot bar above.

```
In [32]: x = [1, 2]
         y = [old_movies, modern_movies]
         label = ['Old Movies', 'Modern Movies']
         plt.figure(figsize = (5,3))
         plt.bar(x, y, tick_label = label, color = 'yellow')
         plt.title('Old and Modern Movies Comparisson by Voting Average')
```

```
plt.xlabel('Type of Movie')
plt.ylabel('Vote Average')
```

```
Out[32]: Text(0,0.5,'Vote Average')
```



It is shown that old movies are more rated than modern movies. Note that the answer to this question is not limited to only these parameters.


1.1.7 3. What is the best month of the year used to release the movie?

- We will start the research of this question plotting the histogram for Release month distribution of movies.

Creating a dataframe containing the release month of each movie to help us reseach this question.

```
In [33]: df_months = df.copy()
df_months['release_month'] = pd.DatetimeIndex(df_months['release_date']).month
df_months['release_month'].hist(figsize = (5, 3))
```

```
Out[33]: <matplotlib.axes._subplots.AxesSubplot at 0x7efc85982978>
```




This shows that most of the movies are released at the end of the year until the beginning of the year (from November to February).

- Plotting the Scatterplot of Release Month in Relation to Revenue.

```
In [34]: df_months.plot(kind = 'scatter', x = 'release_month', y = 'revenue', figsize = (5, 5))
```

```
Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x7efc858961d0>
```



This data don't have any kind of pattern too, which means that no relationship exists between Release month and Revenue.


- To see the bar plot of this question, it will be calculated the mean value of revenue in function of each release month of the movies.

```
In [35]: df_months = df_months.groupby(df_months['release_month'])['revenue'].mean()
```

Having the dataframe for the question defined, there is the plot bar to show the result.

```
In [36]: df_months.plot(kind = 'bar', color = 'green', figsize=(5, 4), alpha = 0.75)
plt.title('Movies Release Month Comparisson by Revenue Value')
plt.xlabel('Release Months')
plt.ylabel('Revenue Value')
```

```
Out[36]: Text(0,0.5,'Revenue Value')
```



The plot shows that movies released in June month tend to be the ones with more revenue value, followed by movies released in May and December months.


Note: These results are not so trustful due to no correlation existing between these two parameters seen in the analysis above.

1.1.8 4. Ranking of the best production companies from year to year in TMDb. By how much?

- Plotting the histogram for Revenue distribution of movies.

```
In [37]: df['revenue'].hist(figsize = (6, 3))
         df['revenue'].median()
```

```
Out[37]: 0.0
```



The plot shows that the distribution is skewed to the right, which means that most of the movies tend to have less than 250 millions of revenue and few of them more than that.

- For the bar plot, it will be used the production companies dataframe created before, where it will be calculated the mean value of the entire dataframe based in each production company in each year.

```
In [38]: p_df = p_df.groupby(['release_year', 'production_companies']).mean()
```

The next step is to get the labels for our plot, that can be determined by maximum revenue value index (that will be in String form containing the release year and production company together, due to the change of index type made in the step above) in each year

```
In [39]: p_df1 = p_df.groupby(['release_year']).idxmax()
labels = p_df1.loc[:, 'revenue']
```

That done, we get the maximum revenue value in each year as height(y) of our plot.

The width(x) or number of bars will correspond to the number of index in one of the dataframes.

```
In [40]: p_df2 = p_df.groupby(['release_year']).max()
y = p_df2.loc[:, 'revenue']

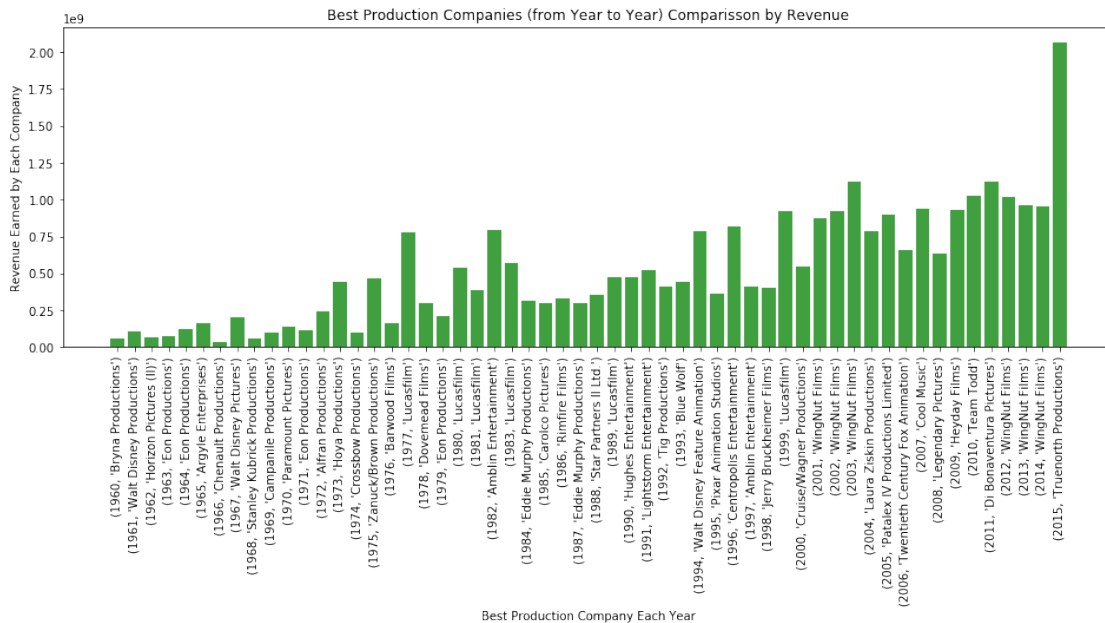
x = p_df2.index.values
```

Finally, the results are above.

```
In [41]: plt.figure(figsize = (16,5))
plt.bar(x, y, tick_label = labels, alpha = 0.75, color = 'green')
```

```
plt.xticks(rotation = 90)
plt.title('Best Production Companies (from Year to Year) Comparisson by Revenue')
plt.xlabel('Best Production Company Each Year')
plt.ylabel('Revenue Earned by Each Company')
```

```
Out[41]: Text(0,0.5,'Revenue Earned by Each Company')
```



Based in this bar plot, the best production company in terms of revenue is the Truenorth Productions that had the highest amount of revenue in 2015 year. Even so, the other production companies shown in the bar plot are the most sucessful in these years too.


Note: These results can be analysed more precisely and accurately with other more parameters too.

1.1.9 5. Does long runtime movies are more expensive to produce compared to short runtime movies?

- We will start the research of this question plotting the histogram for Budget and Runtime Distribution.

```
In [42]: df['budget'].hist(figsize = (5, 3))
         df['budget'].median()
```


```
Out[42]: 200000.0
```

The plot shows that the distribution is skewed to the right, which means that most of the movies tend to have less than 50 millions of budget invested to its production and few of them more than a hundred million.

```
In [43]: df['runtime'].hist(figsize = (5, 3))
         df['runtime'].median()
```

```
Out[43]: 100.0
```




The plot shows that the distribution is skewed to the left, which means that most of the movies tend to have more than 100 runtime minutes and few of them less than 100 minutes.

- Plotting the Scatterplot of Movie Runtime in Relation to Budget.

```
In [44]: df.plot(kind = 'scatter', x = 'runtime', y = 'budget', figsize = (5, 5))
```

```
Out[44]: <matplotlib.axes._subplots.AxesSubplot at 0x7efc85394c88>
```



This data don't have any kind of pattern too, which means that no relationship exists between Runtime and Budget.

- Bar plot for the question research.


It will be defined a median first, which function is to separate short runtime movies from long runtime movies based in the runtime median, so as define the average budget of these two types to compare as next step.

```
In [45]: m = df['runtime'].median()
short_movies = df[df['runtime'] < m]
short_movies = short_movies['budget'].mean()
long_movies = df[df['runtime'] >= m]
long_movies = long_movies['budget'].mean()
```

The answer for the question above.

```
In [46]: x = [1 , 2]
         y = [short_movies, long_movies]
         label = ['Short Movies', 'Long Movies']
         plt.figure(figsize = (5,3))
         plt.bar(x, y, tick_label = label, color = 'yellow')
         plt.title('Long and Short Runtime Movies Comparisson by Budget Needed')
         plt.xlabel('Type of Movie')
         plt.ylabel('Budget')

Out[46]: Text(0,0.5,'Budget')
```




The results above clearly show that long runtime movies can be more expensive to produce.

1.1.10 6. Which genre of movie tends to be more expensive to produce?

To this question, it will be calculated the mean value of budget in function of each genre of movie using the genre dataframe created, as shown above:

```
In [47]: genre_df = g_df.groupby(['genres'])['budget'].mean()
         genre_df.plot(kind = 'bar', color = 'red', figsize=(8, 5))
         plt.title('Movie Genres Comparisson by Budget Average')
         plt.xlabel('Movie Genres')
         plt.ylabel('Budget Average')

Out[47]: Text(0,0.5,'Budget Average')
```




The answer for this question shows that Adventure movies tend to be the most expensive genre to be produced, followed by Fantasy movies.

1.1.11 7. Which movie director directed the production of movies (more than one) that became popular in relation to others, in the most recent year in the dataset?

- Histogram plot of Movie Popularity Distribution

```
In [48]: df['popularity'].hist(figsize = (5, 3))
         df['popularity'].median()
```

```
Out[48]: 0.41976199999999997
```



The plot shows that the distribution of movie popularity is skewed to the right, which means that most of the movies tend to have less than 0.4 of popularity.

- Bar plot for the question research.

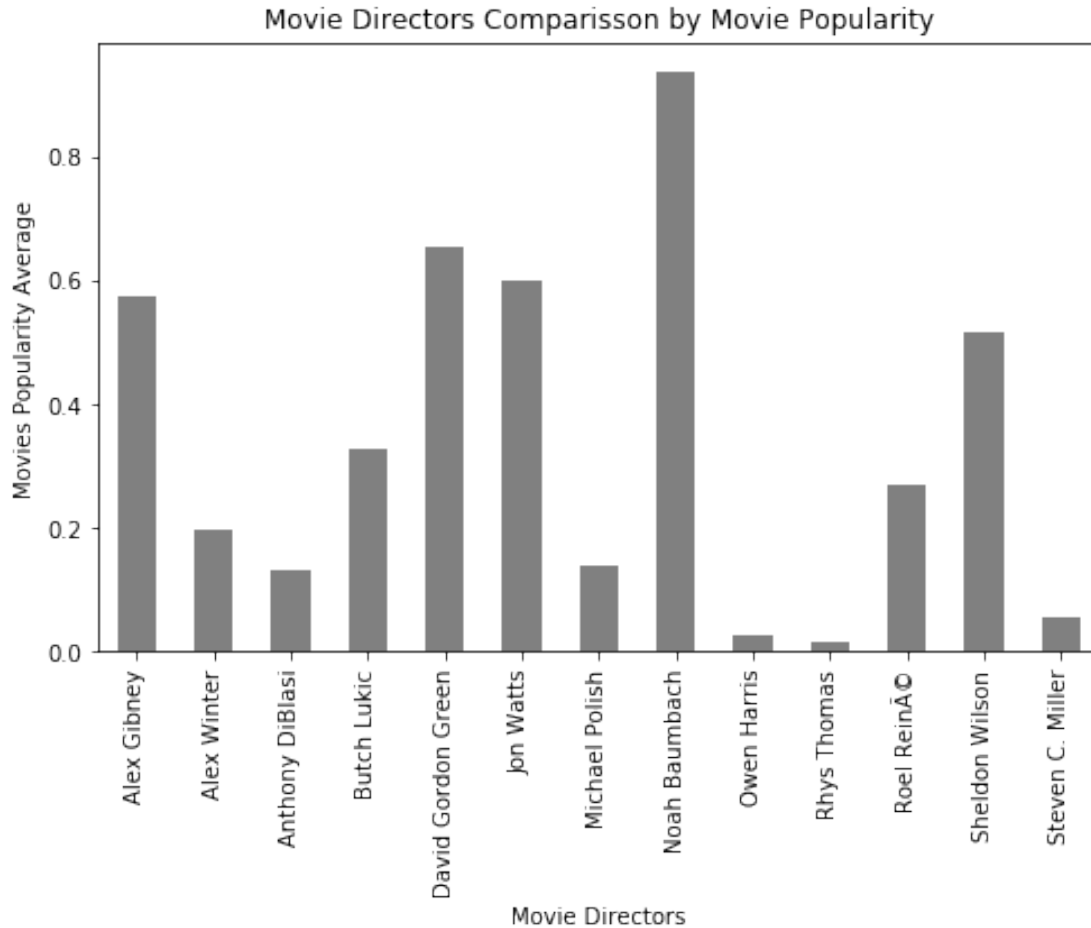
For this question is going to be defined a dataframe that contains the data of movies released in 2015 year and the directors responsables for more than one movies in the same year. That done, it will be calculated the average popularity of the movies in function of each director.

```
In [49]: dir_df = df.query('release_year == 2015')
dir_df = dir_df[dir_df['director'].duplicated()]
dir_df = dir_df.groupby(['director']).mean()['popularity']
```

The answer for the question is above.

```
In [50]: dir_df.plot(kind = 'bar', color = 'gray', figsize=(8, 5))
plt.title('Movie Directors Comparisson by Movie Popularity')
plt.xlabel('Movie Directors')
plt.ylabel('Movies Popularity Average')
```

```
Out[50]: Text(0,0.5,'Movies Popularity Average')
```



The result tells us that Noah Baumbach was the movie director that produced the most popular movies between directors that produced more than one movie in the 2015 year.

Conclusions

As conclusion of this dataset analysis, it will be organized all the findings based in the questions defined above:

The Documentary genre could be the most appreciated movie genre. That said, most of the films tend to have a vote average between 5 to 7 points, and very few of them tend to have more than 8 or less than 4 points as vote, making sense because the Documentary movies genres has the vote average almost to 7 points in relation to other movie genres.

Besides the amount of modern movies is larger than the old movies amount, the old movies had way better vote rating in relation to modern movies, but, it does not mean that there is a relationship between the release year of the movie and the vote average of it, because there are still too many modern movies with good rating points compared to old movies, as it is shown in the scatterplot of the question research.

Most of the movies are released at the end of the year until the beginning of the year (from November to February), but most of the movies released in June month tend to

be the ones with more revenue value, followed by movies released in May and December months. Even so, these results are not so trustful due to no correlation existing between the release month and revenue parameters.

Most of the movies tend to have less than 250 millions of revenue and few of them more than that. That said, the best production company in terms of revenue is the Truenorth Productions that had the highest amount of revenue in 2015 year, which means that it could have been responsible for few movies, but the revenue value was so high that the production company still outstanded.

Most of the movies tend to have less than 50 millions of budget invested to its production and few of them more than a hundred million, so as most of these movies tend to have more than 100 runtime minutes and few of them less than 100 minutes. Regardless the lack of relationship between Runtime and Budget of a movie, long runtime movies can be more expensive to produce. Why? Maybe because these expensive movies are the ones that have more than 100 runtime minutes.

This question can be deeply researched based on other parameters too.

Adventure movies tend to be the most expensive genre to be produced, followed by Fantasy movies.

Lastly, most of the movies trend to have less than 0.4 of popularity, and Noah Baumbach was the movie director that produced the most popular movies between directors that produced more than one movie in the 2015 year (the most recent year included in the dataset).

1.1.12 Limitations

The results can not be so accurate because of the limitation of the number of attributes used, such as genres and production companies of the dataset used to the processes.

The questions were analysed with limited parameters because of the limited scope of analysis, making the results not so accurate and reliable. The statistical result for some questions can be different if we decide to add all the data provided by the dataset.